

# SENTIMENT ANALYSIS-IMDB REVIEWS

BY AAKANKSHA KASHAYAP(21052466) AND RITIKA RANI(21051016)

## import relevant libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
import re
import string
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split as tts
import scipy.stats as stats
import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import roc_curve, roc_auc_score
```

## load the dataset

```
In [2]: data=pd.read_csv('IMDB Dataset.csv')
data
```

Out[2]:

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production.   The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive
...	...	...
49995	I thought this movie did a down right good job...	positive
49996	Bad plot, bad dialogue, bad acting, idiotic di...	negative
49997	I am a Catholic taught in parochial elementary...	negative
49998	I'm going to have to disagree with the previou...	negative
49999	No one expects the Star Trek movies to be high...	negative

50000 rows × 2 columns

## Exploratory Data Analysis

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    review      50000 non-null  object
1    sentiment   50000 non-null  object
dtypes: object(2)
memory usage: 781.4+ KB
```

there is no null values in dataset.

```
In [4]: data.describe()
```

```
Out[4]:
```

	review	sentiment
count	50000	50000
unique	49582	2
top	Loved today's show!!! It was a variety and not...	positive
freq	5	25000

```
In [5]: data['sentiment'].value_counts()
```

```
Out[5]: positive    25000  
negative    25000  
Name: sentiment, dtype: int64
```

dataset is equally divided for positive and negative reviews.

```
In [6]: data.columns
```

```
Out[6]: Index(['review', 'sentiment'], dtype='object')
```

```
In [7]: data.shape
```

```
Out[7]: (50000, 2)
```



```
In [11]: data.head(15)
```

Out[11]:

	review	sentiment
0	one of the other reviewers has mentioned that ...	positive
1	a wonderful little production br br the filmin...	positive
2	i thought this was a wonderful way to spend ti...	positive
3	basically theres a family where a little boy j...	negative
4	petter matteis love in the time of money is a ...	positive
5	probably my alltime favorite movie a story of ...	positive
6	i sure would like to see a resurrection of a u...	positive
7	this show was an amazing fresh innovative ide...	negative
8	encouraged by the positive comments about this...	negative
9	if you like original gut wrenching laughter yo...	positive
10	phil the alien is one of those quirky films wh...	negative
11	i saw this movie when i was about when it cam...	negative
12	so im not a big fan of bolls work but then aga...	negative
13	the cast played shakespearebr br shakespeare l...	negative
14	this a fantastic movie of three prisoners who ...	positive

```
In [12]: def clean_text2(text):
          text=re.sub('[\"'\"\",,]', '', text)
          text=re.sub('\n', '', text)
          return text

          cleaned2=lambda x:clean_text2(x)
```

```
In [13]: data['review']=pd.DataFrame(data.review.apply(cleaned2))
          data.head(15)
```

Out[13]:

	review	sentiment
0	one of the other reviewers has mentioned that ...	positive
1	a wonderful little production br br the filmin...	positive
2	i thought this was a wonderful way to spend ti...	positive
3	basically theres a family where a little boy j...	negative
4	petter matteis love in the time of money is a ...	positive
5	probably my alltime favorite movie a story of ...	positive
6	i sure would like to see a resurrection of a u...	positive
7	this show was an amazing fresh innovative ide...	negative
8	encouraged by the positive comments about this...	negative
9	if you like original gut wrenching laughter yo...	positive
10	phil the alien is one of those quirky films wh...	negative
11	i saw this movie when i was about when it cam...	negative
12	so im not a big fan of bolls work but then aga...	negative
13	the cast played shakespearebr br shakespeare l...	negative
14	this a fantastic movie of three prisoners who ...	positive

```
In [14]: x_review = data.iloc[0:,0].values
          y_sentiment = data.iloc[0:,1].values
```

```
In [15]: x_train,x_test,y_train,y_test = tts(x_review,y_sentiment,test_size = 0.20,random_state = 41)
```


## Extracting the data

```
In [16]: > tf = TfidfVectorizer()  
from sklearn.pipeline import Pipeline
```

```
In [17]: > label_encoder = LabelEncoder()  
y_train_encoded = label_encoder.fit_transform(y_train)  
y_test_encoded = label_encoder.transform(y_test)
```

## Modeling-Logistic regression

```
In [18]: > from sklearn.linear_model import LogisticRegression  
classifier=LogisticRegression()  
model_logistic=Pipeline([('vectorizer',tf),('classifier',classifier)])  
model_logistic.fit(x_train,y_train_encoded)
```



```
In [19]: > y_pred=model_logistic.predict(x_test)
```

```
In [20]: > # model score  
accuracy_score(y_pred,y_test_encoded)
```

Out[20]: 0.8977

```
In [21]: > conf_mat=confusion_matrix(y_test_encoded,y_pred)  
print(conf_mat)
```


```
[[4421  559]  
 [ 464 4556]]
```

```
In [22]: > # f1 score  
recall=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[1][0])  
precision=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[0][1])  
F1=2*recall*precision/(recall+precision)  
print(recall)  
print(precision)  
print(F1)
```

```
0.9050153531218015  
0.8877510040160642  
0.8963000506842371
```

## Modeling-Decision Tree

```
In [23]: > from sklearn.tree import DecisionTreeClassifier  
from sklearn.pipeline import Pipeline  
  
classifier=DecisionTreeClassifier(max_depth=20,min_samples_split=10,min_samples_leaf=5)  
model_decisiontree = Pipeline([('vectorizer', TfidfVectorizer(max_features=10000)), ('classifier', clas  
model_decisiontree.fit(x_train, y_train_encoded)
```



```
In [24]: y_pred=model_decisiontree.predict(x_test)
```

```
In [25]: accuracy_score(y_pred,y_test_encoded)
```

```
Out[25]: 0.7351
```

```
In [26]: conf_mat=confusion_matrix(y_test_encoded,y_pred)
print(conf_mat)
```

```
[[3295 1685]
 [ 964 4056]]
```

```
In [27]: recall=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[1][0])
precision=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[0][1])
F1=2*recall*precision/(recall+precision)
print(recall)
print(precision)
print(F1)
```

```
0.7736557877436018
0.6616465863453815
0.7132806580798787
```

## Modeling-Linear Regression

```
In [28]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

regressor = LinearRegression()
model_linear= Pipeline([('vectorizer', TfidfVectorizer()), ('regressor', regressor)])
model_linear.fit(x_train, y_train_encoded)

y_pred = model_linear.predict(x_test)
mse = mean_squared_error(y_test_encoded, y_pred)
print("Mean Squared Error:", mse)
```

```
Mean Squared Error: 0.1828266194890447
```

## Modeling-Naive Bayes

```
In [29]: from sklearn.naive_bayes import MultinomialNB
classifier = MultinomialNB()

model_naivebayes=Pipeline([('vectorizer',TfidfVectorizer()),('classifier',classifier)])
model_naivebayes.fit(x_train, y_train_encoded)
y_pred=model_naivebayes.predict(x_test)
accuracy_score(y_pred,y_test_encoded)
```

```
Out[29]: 0.8615
```

```
In [30]: conf_mat=confusion_matrix(y_test_encoded,y_pred)
print(conf_mat)
```

```
[[4423  557]
 [ 828 4192]]
```

```
In [31]: recall=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[1][0])
precision=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[0][1])
F1=2*recall*precision/(recall+precision)
print(recall)
print(precision)
print(F1)
```

```
0.8423157493810702
0.8881526104417671
0.8646271136741276
```

## Modeling-KNN

```
In [32]: from sklearn.neighbors import KNeighborsClassifier

classifier=KNeighborsClassifier(n_neighbors=5, algorithm='auto', n_jobs=-1)
model_knn=Pipeline([('vectorizer',TfidfVectorizer(max_features=10000)),('classifier', classifier)])
model_knn.fit(x_train, y_train_encoded)

y_pred=model_knn.predict(x_test)
accuracy_score(y_pred,y_test_encoded)
```

Out[32]: 0.7465

```
In [33]: conf_mat=confusion_matrix(y_test_encoded,y_pred)
print(conf_mat)
```

```
[[3511 1469]
 [1066 3954]]
```

```
In [34]: recall=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[1][0])
precision=conf_mat[0][0]/(conf_mat[0][0]+conf_mat[0][1])
F1=2*recall*precision/(recall+precision)
print(recall)
print(precision)
print(F1)
```

```
0.7670963513218265
0.7050200803212852
0.7347493983467615
```

## Result

```
In [35]: metrics_df={'Algorithms':['Logistic','knn','naive bayes','decision tree'],
                    'Accuracy':[89.77,74.65,86.15,73.42],
                    'Precision':[88.77,70.50,88.81,65.76],
                    'Recall':[90.5,76.7,84.23,77.46],
                    'F1-Score':[89.63,73.47,86.46,71.13]}
df=pd.DataFrame(metrics_df)
df
```

Out[35]:

	Algorithms	Accuracy	Precision	Recall	F1-Score
0	Logistic	89.77	88.77	90.50	89.63
1	knn	74.65	70.50	76.70	73.47
2	naive bayes	86.15	88.81	84.23	86.46
3	decision tree	73.42	65.76	77.46	71.13

linear regression algorithm does not support f1 score, precision, recall, and confusion matrix for categorical variables. we have shown calculation for mse of linear regression algorithm.  
confusion matrix for every algorithm is also calculated.

## Plotting the ROC curve

```

In [37]: from sklearn.metrics import roc_curve, auc

y_prob_logistic = model_logistic.predict_proba(x_test)[: , 1]
y_prob_tree = model_decisiontree.predict_proba(x_test)[: , 1]
y_prob_linear = model_linear.predict(x_test) # No predict_proba for Linear Regression
y_prob_nb = model_naivebayes.predict_proba(x_test)[: , 1]
y_prob_knn = model_knn.predict_proba(x_test)[: , 1]

# Compute ROC curves and ROC areas for the positive class
fpr_logistic, tpr_logistic, thresholds_logistic = roc_curve(y_test_encoded, y_prob_logistic)
roc_auc_logistic = auc(fpr_logistic, tpr_logistic)

fpr_tree, tpr_tree, thresholds_tree = roc_curve(y_test_encoded, y_prob_tree)
roc_auc_tree = auc(fpr_tree, tpr_tree)

fpr_linear, tpr_linear, thresholds_linear = roc_curve(y_test_encoded, y_prob_linear)
roc_auc_linear = auc(fpr_linear, tpr_linear)


fpr_nb, tpr_nb, thresholds_nb = roc_curve(y_test_encoded, y_prob_nb)
roc_auc_nb = auc(fpr_nb, tpr_nb)

fpr_knn, tpr_knn, thresholds_knn = roc_curve(y_test_encoded, y_prob_knn)
roc_auc_knn = auc(fpr_knn, tpr_knn)

# Plot ROC curves for all models
plt.figure(figsize=(10, 8))
plt.plot(fpr_logistic, tpr_logistic, color='darkorange', lw=2, label=f'Logistic Regression(AUC={roc_auc_log:.2f})')
plt.plot(fpr_tree, tpr_tree, color='green', lw=2, label=f'Decision Tree(AUC={roc_auc_tree:.2f})')
plt.plot(fpr_linear, tpr_linear, color='blue', lw=2, label=f'Linear Regression(AUC={roc_auc_linear:.2f})')
plt.plot(fpr_nb, tpr_nb, color='purple', lw=2, label=f'Naive Bayes(AUC={roc_auc_nb:.2f})')
plt.plot(fpr_knn, tpr_knn, color='red', lw=2, label=f'KNN(AUC={roc_auc_knn:.2f})')

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curves')
plt.legend(loc='lower right')
plt.show()

```





## conclusion

1. for sentiment analysis, logistic regression has best result than other algorithms.
2. we have shown f1 score, recall, precision and confusion matrix for algorithms-knn,naive bayes,logistic regression,linear regression,decision tree.
3. we have plot roc curve for all above algorithms.
4. wordcloud is used to show common and different words for positive and negative reviews.